

Data Structures: Vocab, Lexemes and StringStore

ADVANCED NLP WITH SPACY



Ines Montani
spaCy core developer

Shared vocab and string store (1)

- `Vocab` : stores data shared across multiple documents
- To save memory, spaCy encodes all strings to **hash values**
- Strings are only stored once in the `StringStore` via `nlp.vocab.strings`
- String store: **lookup table** in both directions

```
coffee_hash = nlp.vocab.strings['coffee']  
coffee_string = nlp.vocab.strings[coffee_hash]
```

- Hashes can't be reversed – that's why we need to provide the shared vocab

```
# Raises an error if we haven't seen the string before  
string = nlp.vocab.strings[3197928453018144401]
```

Shared vocab and string store (2)

- Look up the string and hash in `nlp.vocab.strings`

```
doc = nlp("I love coffee")
print('hash value:', nlp.vocab.strings['coffee'])
print('string value:', nlp.vocab.strings[3197928453018144401])
```

```
hash value: 3197928453018144401
string value: coffee
```

- The `doc` also exposes the vocab and strings

```
doc = nlp("I love coffee")
print('hash value:', doc.vocab.strings['coffee'])
```

```
hash value: 3197928453018144401
```

Lexemes: entries in the vocabulary

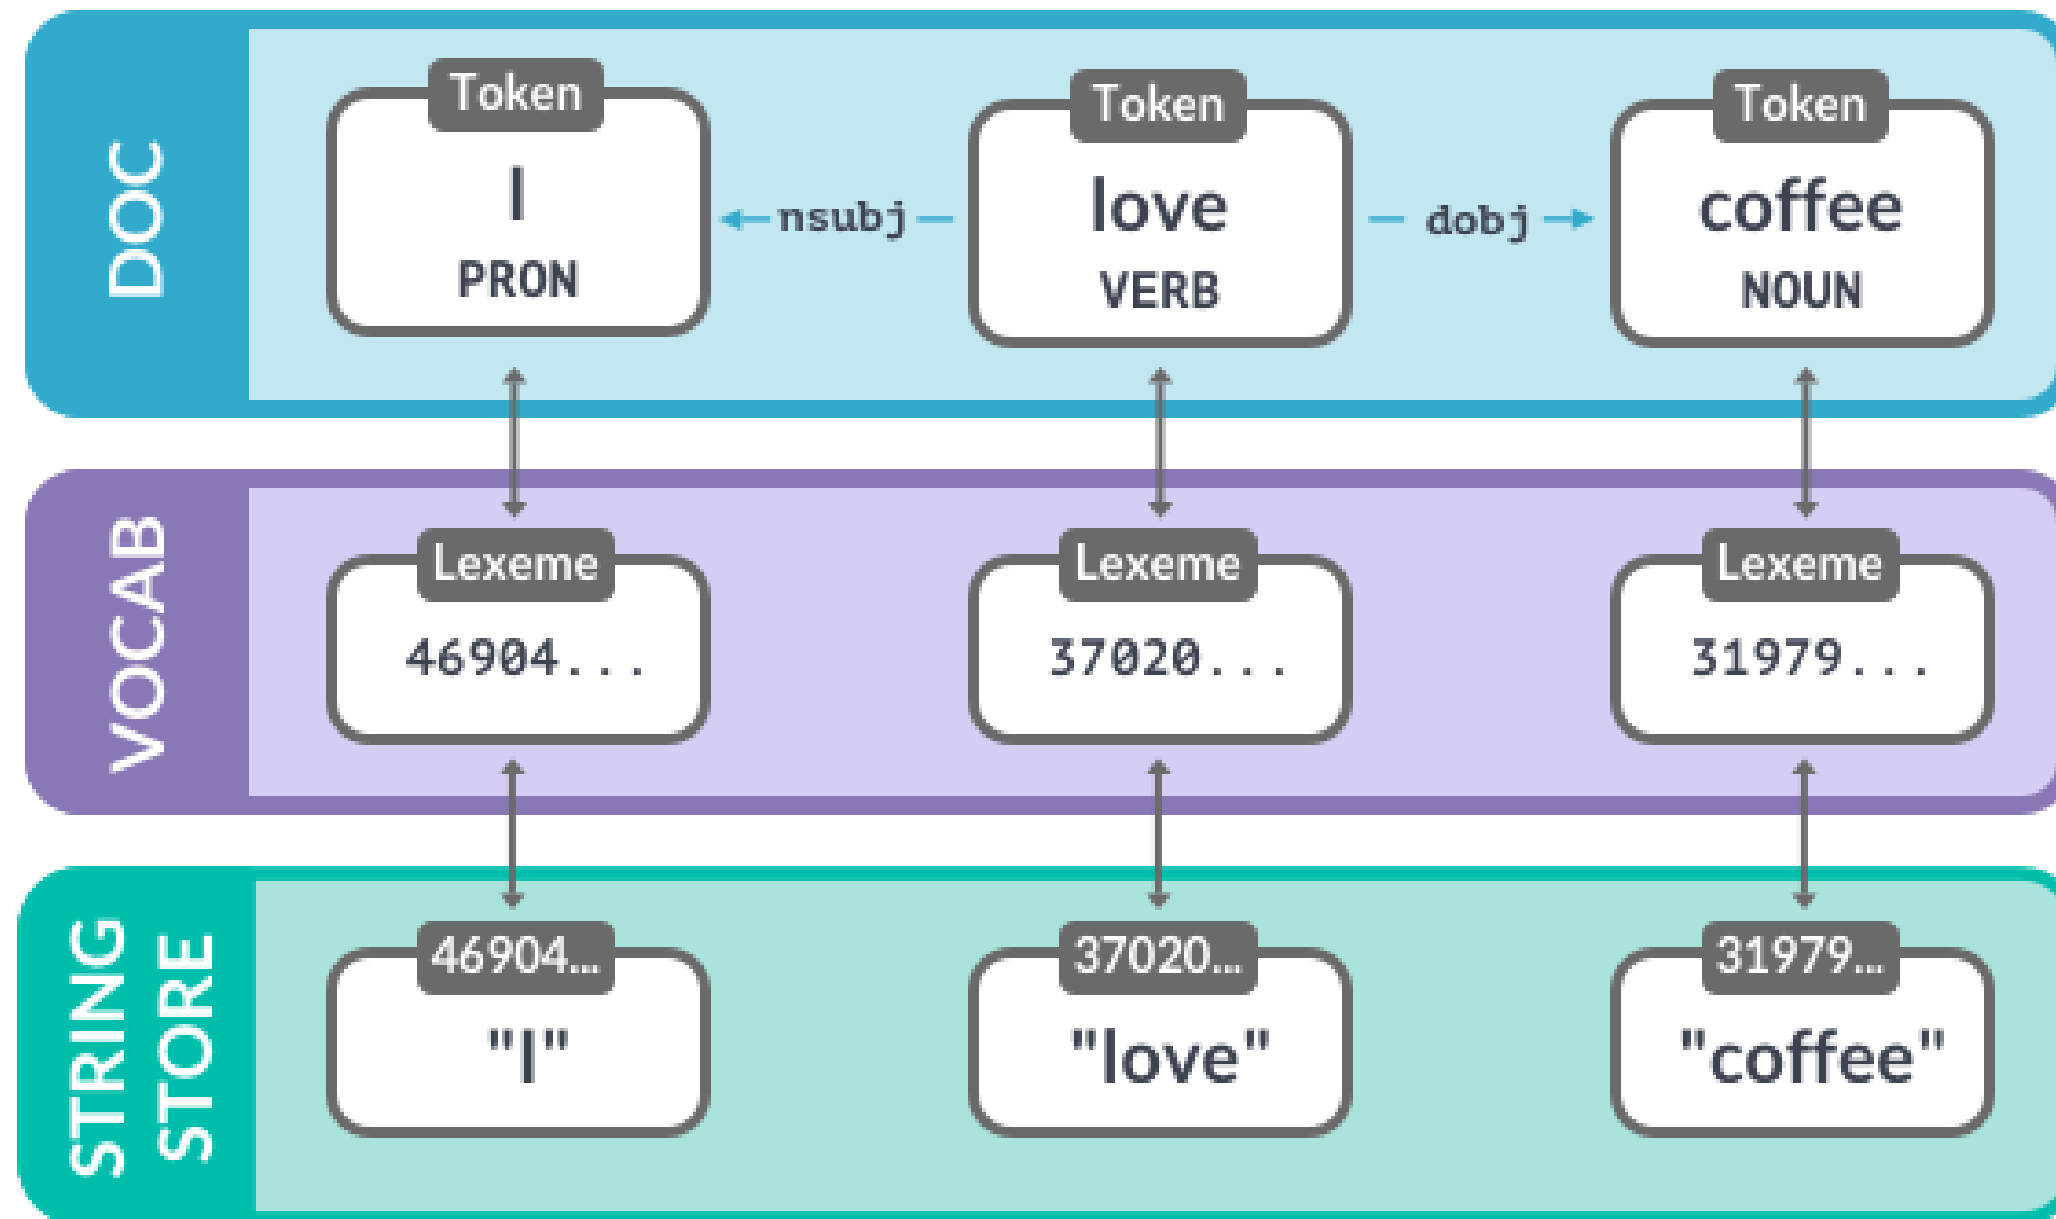
- A `Lexeme` object is an entry in the vocabulary

```
doc = nlp("I love coffee")
lexeme = nlp.vocab['coffee']
# print the lexical attributes
print(lexeme.text, lexeme.orth, lexeme.is_alpha)
```

```
coffee 3197928453018144401 True
```

- Contains the **context-independent** information about a word
 - Word text: `lexeme.text` and `lexeme.orth` (the hash)
 - Lexical attributes like `lexeme.is_alpha`
 - **Not** context-dependent part-of-speech tags, dependencies or entity labels


Vocab, hashes and lexemes



Let's practice!
ADVANCED NLP WITH SPACY

Data Structures: Doc, Span and Token

ADVANCED NLP WITH SPACY



Ines Montani
spaCy core developer

The Doc object

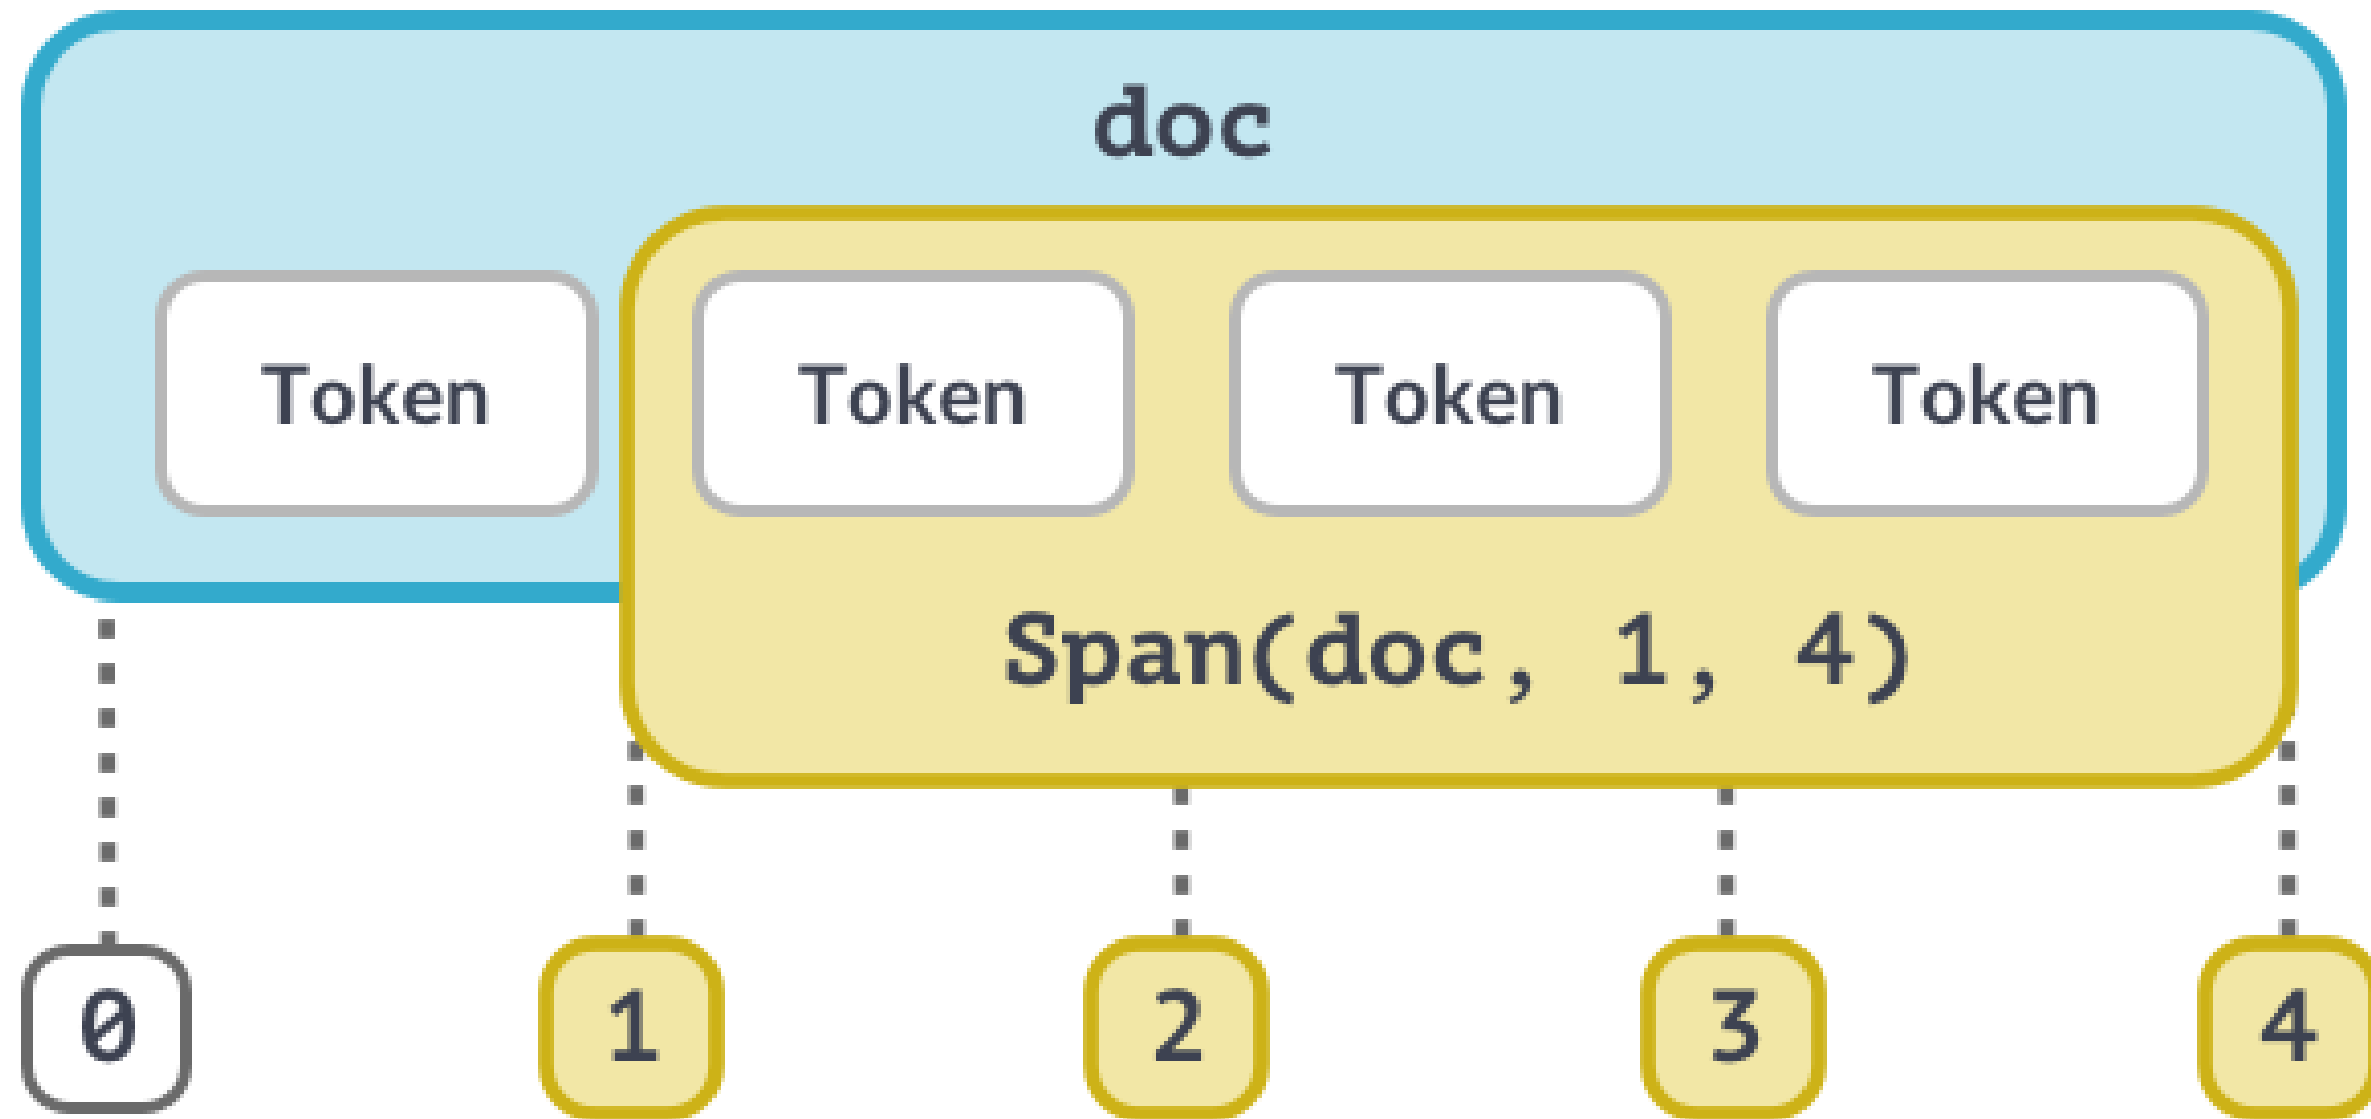
```
# Create an nlp object
from spacy.lang.en import English
nlp = English()

# Import the Doc class
from spacy.tokens import Doc

# The words and spaces to create the doc from
words = ['Hello', 'world', '!']
spaces = [True, False, False]

# Create a doc manually
doc = Doc(nlp.vocab, words=words, spaces=spaces)
```


The Span object (1)



The Span object (2)

```
# Import the Doc and Span classes
from spacy.tokens import Doc, Span

# The words and spaces to create the doc from
words = ['Hello', 'world', '!']
spaces = [True, False, False]

# Create a doc manually
doc = Doc(nlp.vocab, words=words, spaces=spaces)

# Create a span manually
span = Span(doc, 0, 2)

# Create a span with a label
span_with_label = Span(doc, 0, 2, label="GREETING")

# Add span to the doc.ents
doc.ents = [span_with_label]
```


Best practices

- `Doc` and `Span` are very powerful and hold references and relationships of words and sentences
 - **Convert result to strings as late as possible**
 - **Use token attributes if available** – for example, `token.i` for the token index
- Don't forget to pass in the shared `vocab`

Let's practice!
ADVANCED NLP WITH SPACY

Word vectors and semantic similarity

ADVANCED NLP WITH SPACY



Ines Montani

spaCy core developer

Comparing semantic similarity

- `spacy` can compare two objects and predict similarity
- `Doc.similarity()` , `Span.similarity()` and `Token.similarity()`
- Take another object and return a similarity score (`0` to `1`)
- **Important:** needs a model that has word vectors included, for example:
 - **YES:** `en_core_web_md` (medium model)
 - **YES:** `en_core_web_lg` (large model)
 - **NO:** `en_core_web_sm` (small model)

Similarity examples (1)

```
# Load a larger model with vectors
nlp = spacy.load('en_core_web_md')
# Compare two documents
doc1 = nlp("I like fast food")
doc2 = nlp("I like pizza")
print(doc1.similarity(doc2))
```

0.8627204117787385

```
# Compare two tokens
doc = nlp("I like pizza and pasta")
token1 = doc[2]
token2 = doc[4]
print(token1.similarity(token2))
```

0.7369546

Similarity examples (2)

```
# Compare a document with a token
```

```
doc = nlp("I like pizza")
```

```
token = nlp("soap")[0]
```

```
print(doc.similarity(token))
```

```
0.32531983166759537
```

```
# Compare a span with a document
```

```
span = nlp("I like pizza and pasta")[2:5]
```

```
doc = nlp("McDonalds sells burgers")
```

```
print(span.similarity(doc))
```

```
0.619909235817623
```


How does spaCy predict similarity?

- Similarity is determined using **word vectors**
- Multi-dimensional meaning representations of words
- Generated using an algorithm like **Word2Vec** and lots of text
- Can be added to spaCy's statistical models
- Default: cosine similarity, but can be adjusted
- **Doc** and **Span** vectors default to average of token vectors
- Short phrases are better than long documents with many irrelevant words

Word vectors in spaCy

```
# Load a larger model with vectors
nlp = spacy.load('en_core_web_md')
doc = nlp("I have a banana")
# Access the vector via the token.vector attribute
print(doc[3].vector)
```

```
[2.02280000e-01, -7.66180009e-02,  3.70319992e-01,
 3.28450017e-02, -4.19569999e-01,  7.20689967e-02,
-3.74760002e-01,  5.74599989e-02, -1.24009997e-02,
 5.29489994e-01, -5.23800015e-01, -1.97710007e-01,
-3.41470003e-01,  5.33169985e-01, -2.53309999e-02,
 1.73800007e-01,  1.67720005e-01,  8.39839995e-01,
 5.51070012e-02,  1.05470002e-01,  3.78719985e-01,
 2.42750004e-01,  1.47449998e-02,  5.59509993e-01,
 1.25210002e-01, -6.75960004e-01,  3.58420014e-01,
-4.00279984e-02,  9.59490016e-02, -5.06900012e-01,
-8.53179991e-02,  1.79800004e-01,  3.38669986e-01,
...]
```

Similarity depends on the application context

- Useful for many applications: recommendation systems, flagging duplicates etc.
- There's no objective definition of "similarity"
- Depends on the context and what application needs to do

```
doc1 = nlp("I like cats")  
doc2 = nlp("I hate cats")  
  
print(doc1.similarity(doc2))
```

```
0.9501447503553421
```

Let's practice!
ADVANCED NLP WITH SPACY

Combining models and rules

ADVANCED NLP WITH SPACY



Ines Montani

spaCy core developer

Statistical predictions vs. rules

	Statistical models	Rule-based systems
Use cases	application needs to <i>generalize</i> based on examples	
Real-world examples	product names, person names, subject/object relationships	
spaCy features	entity recognizer, dependency parser, part-of-speech tagger	

Statistical predictions vs. rules

	Statistical models	Rule-based systems
Use cases	application needs to <i>generalize</i> based on examples	dictionary with finite number of examples
Real-world examples	product names, person names, subject/object relationships	countries of the world, cities, drug names, dog breeds
spaCy features	entity recognizer, dependency parser, part-of-speech tagger	tokenizer, <code>Matcher</code> , <code>PhraseMatcher</code>

Recap: Rule-based Matching

```
# Initialize with the shared vocab
from spacy.matcher import Matcher
matcher = Matcher(nlp.vocab)

# Patterns are lists of dictionaries describing the tokens
pattern = [{'LEMMA': 'love', 'POS': 'VERB'}, {'LOWER': 'cats'}]
matcher.add('LOVE_CATS', None, pattern)

# Operators can specify how often a token should be matched
pattern = [{'TEXT': 'very', 'OP': '+'}, {'TEXT': 'happy'}]

# Calling matcher on doc returns list of (match_id, start, end) tuples
doc = nlp("I love cats and I'm very very happy")
matches = matcher(doc)
```


Adding statistical predictions

```
matcher = Matcher(nlp.vocab)
matcher.add('DOG', None, [{'LOWER': 'golden'}, {'LOWER': 'retriever'}])
doc = nlp("I have a Golden Retriever")
for match_id, start, end in matcher(doc):
    span = doc[start:end]
    print('Matched span:', span.text)
    # Get the span's root token and root head token
    print('Root token:', span.root.text)
    print('Root head token:', span.root.head.text)
    # Get the previous token and its POS tag
    print('Previous token:', doc[start - 1].text, doc[start - 1].pos_)
```

```
Matched span: Golden Retriever
Root token: Retriever
Root head token: have
Previous token: a DET
```

Efficient phrase matching (1)

- `PhraseMatcher` like regular expressions or keyword search – but with access to the tokens!
- Takes `Doc` object as patterns
- More efficient and faster than the `Matcher`
- Great for matching large word lists

Efficient phrase matching (2)

```
from spacy.matcher import PhraseMatcher

matcher = PhraseMatcher(nlp.vocab)
pattern = nlp("Golden Retriever")
matcher.add('DOG', None, pattern)

doc = nlp("I have a Golden Retriever")

# iterate over the matches
for match_id, start, end in matcher(doc):
    # get the matched span
    span = doc[start:end]
    print('Matched span:', span.text)
```

Matched span: Golden Retriever

Let's practice!
ADVANCED NLP WITH SPACY